

# Database Theory

William Schultz

November 8, 2022

## Joins

At a high level, database (i.e. SQL) tables can be viewed as  $n$ -ary relations (or, more plainly, as “spreadsheets”). For example, consider the following relation  $P$

Name	Age	HouseId
Alice	31	6
Bob	32	3
Jane	25	4

and another relation  $H$

Id	Year
6	1904
4	1965

We can consider the *cross product* of these two relations  $P \times H$ , which is simply the Cartesian product of all rows (i.e. tuples) in  $P$  with all rows in  $H$ , giving relation  $P \times H$  as

Name	Age	HouseId	Id	Year
Alice	31	6	6	1904
Alice	31	6	4	1965
Bob	32	3	6	1904
Bob	32	3	4	1965
Jane	25	4	6	1904
Jane	25	4	4	1965

with a total tuple count of  $|P \times H| = |P| \times |H| = 6$ .

On its own, the full cross product of two tables may not be very useful, but a commonly useful operation to apply after doing this cross product is the *join*, which essentially just applies some filter (e.g. predicate) to the tuples that are generated as a result of this cross product operation. If we filter the result  $P \times H$  based on the predicate  $HouseId == Id$ , then we say we’re “joining” the two tables,  $P$  and  $H$ , on  $HouseId == Id$ , which gives as a result:

Name	Age	HouseId	Id	Year
Alice	31	6	6	1904
Jane	25	4	4	1965

which is basically the set of people in  $P$  associated with the house in  $H$  they own. More compactly, we typically notate a join on predicate  $p$  between two relations  $A$  and  $B$  as

$$A \bowtie_p B$$

Again, we can think of this as simply a composition of cross product and filtering operations i.e.

$$A \bowtie_p B = \sigma_p(A \times B)$$

where  $\sigma_p$  represents the filtering operation for a given predicate  $p$  on tuples.

Note that joins are *commutative* i.e.  $A \bowtie_p B = B \bowtie_p A$ . This can be easily seen from examining the decomposed form of joins in terms of cross products and filtering i.e. since  $A \times B = B \times A$  (if we ignore ordering of columns in the output tuples). Note also that we can view a sequence of joins as a big cross product followed by a filtering operation at the end i.e.

$$(A \bowtie_p B) \bowtie_q C = \sigma_q(\sigma_p(A \times B) \times C) = \sigma_{p \wedge q}(A \times B \times C)$$

Furthermore, joins are *associative*. This means that we can re-order joins as we please (since they are both *associative* and *commutative*), so there may be many different join executing orderings that produce the same final result (TODO: query optimization). That is,

$$\begin{aligned}
& (A \bowtie_p B) \bowtie_q C \\
& A \bowtie_p (B \bowtie_q C) \\
& A \bowtie_p (C \bowtie_q B) \\
& (A \bowtie_p C) \bowtie_q B \\
& (C \bowtie_p A) \bowtie_q B
\end{aligned}$$

are all equivalent. Some orderings may, however, be much more efficient to execute. More generally, we can think of a particular join ordering as a binary tree, that essentially maps to the syntactic parse tree of the above expressions.